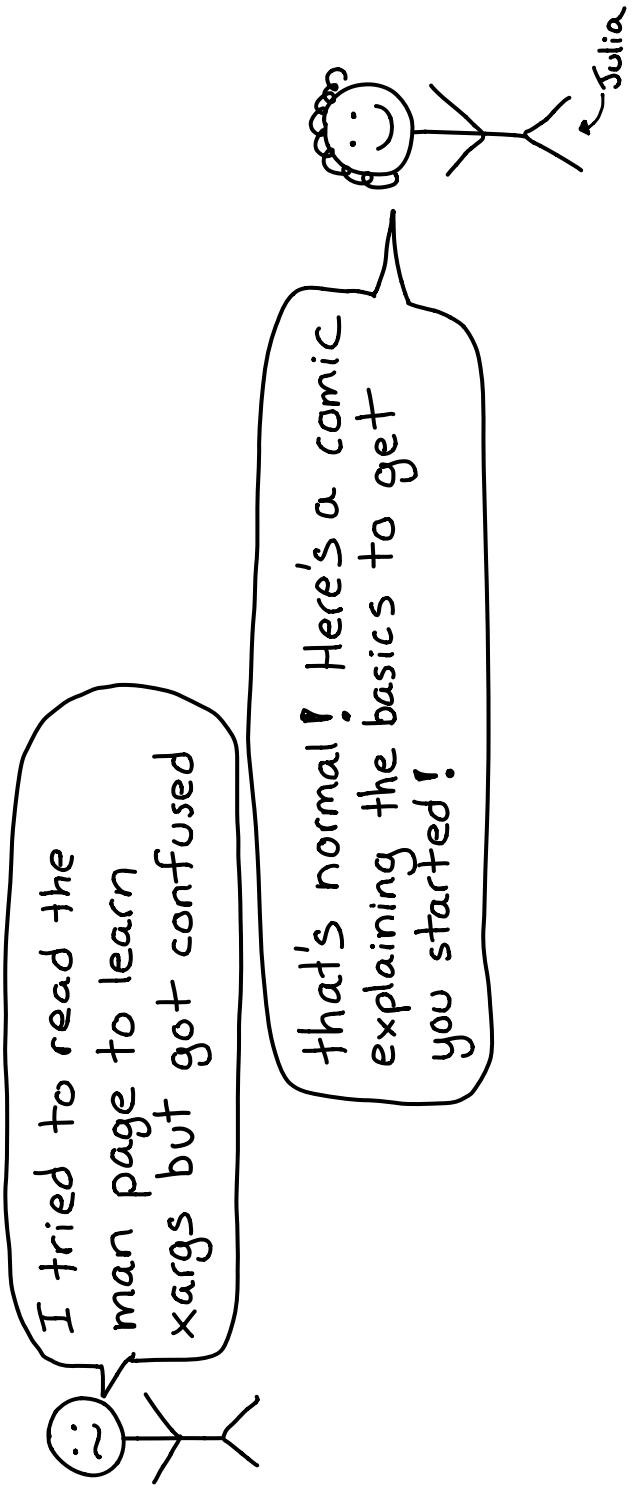




love this?
find more awesome zines at
→ wizardzines.com ←

This zine explains some of the most useful Unix command line tools in 1 page each.



Even if you've used the tool before, I might have a new trick or two for you ♡

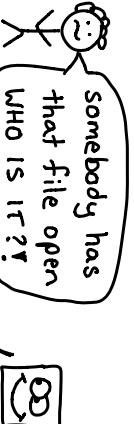
more useful tools

- make
- jq
- nohup
- disown
- cut/paste
- sponge
- xxd
- hexdump
- objdump
- strings
- screen
- tmux
- date
- entr
- seq
- join
- parallel :
- GNU parallel
- pigz / pixz
- sort --parallel
- diff -U
- vipe
- imagemagick
- fish
- ranger
- chronic

lsof

lsof

stands for list open files



I can tell you!

-P PID

for each open file:
→ pid

- file type (regular? directory?
FIFO? socket?)
- file descriptor (FD column)
- user
- filename / socket address

list the files PID has
open

lsof /some /dir
list just the open files
in /some /dir

-i

list open network sockets
(sockets are files!)

examples:

- i -n -P ← -n & -P mean "don't resolve host names / ports"
- i TCP (also -Pni)
- i -s TCP:LISTEN

find deleted files

\$ lsof | grep deleted
will show you deleted files!

You can recover open deleted
files from
/proc/PID/fd/FD
process that opened the file

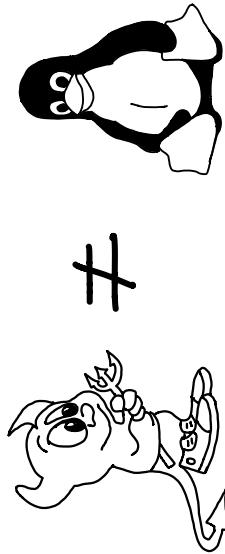
netstat

another way to list open
sockets on Linux is:

netstat -tunapl
↑
+unar, please!
On Mac, netstat has
different args.

♥ Table of contents ♥

BSD≠GNU.....4	bash tricks.....10-11	misc commands...17
grep.....5	disk usage.....12	head & tail.....18
find.....6	tar.....13	less.....19
xargs.....7	ps.....14	kill.....20
awk.....8	top.....15	cat.....21
sed.....9	sort & uniq.....16	lsof.....22

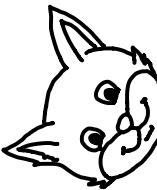


For almost all of these tools, there are at least 2 versions:

① The BSD version (on BSDs & Mac OS)

② The GNU version (on Linux)

All of the examples in this zine were tested on Linux. Some things like sed -i are different on Mac. Be careful when writing cross-platform scripts! You can install the GNU versions on Mac with brew install coreutils.



cat & friends

21

cat concatenates files

```
$ cat myfile.txt
prints contents of myfile.txt
$ cat *.txt
prints all .txt files put together!
```

you can use cat as an EXTREMELY BASIC text editor:

- ① Run \$ cat > file.txt
- ② type the contents (don't make mistakes !)
- ③ press ctrl+d to finish

cat -n

prints out the file with line numbers!

- 1 Once upon a midnight...
- 2 Over many a quaint...
- 3 While I nodded, nearly

zcat

cats a gzipped file!
Actually just a 1-line shell script that runs gzip -cd, but easier to remember.

tee

'tee file.txt' will write its stdin to both stdout and file.txt
stdin → tee a.txt → stdout → a.txt

```
how to redirect to
a file owned by root
$ sudo echo "hi" >> x.txt
this will open x.txt as your
user, not as root, so it fails!
$ echo "hi" | sudo tee -a x.txt
will open x.txt as root ``
```

Kill

kill doesn't just kill programs

 You can send ANY signal to a program with kill!

\$ kill -SIGNAL PID

name or number

killall -SIGNAL NAME

signals all processes called NAME for example:

\$ killall firefox

useful flags:

- w wait for all signaled processes to die
- i ask before signalling

	name	num
kill	=>	SIGTERM 15
kill -9	=>	SIGKILL 9
kill -KILL	=>	SIGKILL 9
kill -HUP	=>	SIGHUP 1
kill -STOP	=>	SIGSTOP 19

1 HUP 2 INT 3 QUIT 4 ILL
 5 TRAP 6 ABRT 7 BUS 8 FPE
 9 KILL 10 USR1 11 SEGV 12 USR2
 13 PIPE 14 ALRM 15 TERM 16 SIGFLT
 17 CHLD 18 CONT 19 STOP 20 TSTP
 21 TTIN 22 TTOUT 23 URG 24 XCPU
 25 XFSZ 26 VTIME 27 PROF 28 WINCH
 29 POLL 30 PUR 31 SIS

	lists all signals
kill -l	1 HUP 2 INT 3 QUIT 4 ILL 5 TRAP 6 ABRT 7 BUS 8 FPE 9 KILL 10 USR1 11 SEGV 12 USR2 13 PIPE 14 ALRM 15 TERM 16 SIGFLT 17 CHLD 18 CONT 19 STOP 20 TSTP 21 TTIN 22 TTOUT 23 URG 24 XCPU 25 XFSZ 26 VTIME 27 PROF 28 WINCH 29 POLL 30 PUR 31 SIS

pgrep

prints PIDs of matching running programs

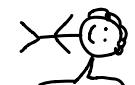
pgrep fire matches firefox firebird

NOT bash firefox.sh

To search the whole command line (eg bash firefox.sh), use `pgrep -f`

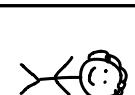
same as pgrep, but signals PIDs found. Example:

\$ pkill -f firefox

 I use pkill more than killall these days

pkill

recursive! Search all the files in a directory.

 only print the matching part of the line instead of the whole line

Grep

5

use if you want regexps like ".+" to work. otherwise you need to use ".\+" aka egrep

 recursive! Search all the files in a directory.

 invert match: find all lines that don't match

 only show the filenames of the files that matched

 case insensitive

 only show the string as a regex

 Show context for your search. For example:

 \$ grep -A 3 cats will show 3 lines of context after a match

 don't treat the match

 aka fgrep

 grep alternatives

 ack, ag, ripgrep (better for searching code!) ripgrep

find

6

find searches a directory for files
\$ find /tmp -type d -print
↑
directory which files action to do with the files
to search
 here are my favourite find arguments!

-name/-iname
case insensitive
the filename! Example:
-name '*.txt'

-path / -i-path
search the full path!
-path '/home/*/*.go'

-type TYPE
f: regular file l: symlink
d: directory + more!

-maxdepth NUM
only descend NUM levels when searching a directory

-size 0
find empty files!
Useful to find files you created by accident

-exec COMMAND
action: run COMMAND on every file found

-print0
print null-separated filenames.
Use with xargs -0!

-delete
action: delete all files found

less

19

less is a pager
that means it lets you view (not edit) text files or piped-in text
man uses your pager (usually less) to display man pages

many vim shortcuts work in less
/ search
n/N next / prev match
j/k down / up a line
m/ mark / return to line
g/G beginning / end of file
(gg in vim)

less -r
displays bash escape codes as colours
try \$ ls --color | less -r
with -r without -r
a.txt → ESC[31m a.txt
a.txt.gz → ESC[31m a.txt.gz
red, bold

q quit ↵
V lowercase edit file in your \$EDITOR
arrow keys, Home / End, Pg Up, Pg Dn work in less

F press F to keep reading from the file as it's updated (like tail -f)
press Ctrl+C to stop reading updates

+ runs a command when less starts
less +F : follow updates
less +G : start at end of file
less +20% : start 20% into file
less +/foo : search for 'foo' right away

head & tail

head
shows you the first 10 lines of a file

if you pipe a program's output to head, the program will stop after printing 10 lines (it gets sent ~~SIGPIPE~~)

-c NUM

show the first/last NUM bytes of the file

\$ head -c 1k

will show the first 1024 bytes

tail
tail shows the last 10 lines!
{tail -f FILE} will follow:

print any new lines added to the end of FILE. Super useful for log files!

tail --retry

keep trying to open file if it's inaccessible

tail --pid PID

stop when process PID stops running (with -f)

-n NUM
-n NUM (either head or tail) will change the # lines shown

NUM can also be negative. Example:
\$ head -n -5 file.txt
will print all lines except the last 5

tail --follow=name

Usually tail -f will follow a file descriptor.

tail --follow=name FILENAME will keep following the same filename, even if the file descriptor changes

xargs

7

xargs takes whitespace separated strings from stdin and converts them into command-line arguments

echo /home/tmp | xargs ls

will run:

ls /home/tmp

this is useful when you want to run the same command on a list of files!

- delete (xargs rm)
- combine (xargs cat)
- search (xargs grep)
- replace (xargs sed)

how to lint every Python file in your Git repo:

git ls-files | grep .py | xargs pep8

if there are spaces in your filenames "my day.txt"

xargs will think it's 2 files "my" and "day.txt"

fix it like this:

find . -print0 | xargs -0 COMMAND

more useful xargs options:

(-n 1) makes xargs run max-args for each input capital P is the max number of parallel processes xargs will start

awk

8

awk is a tiny programming language for manipulating columns of data

I only know how to do 2 things with awk but it's still useful!

basic awk program structure:

```
BEGIN { ... }
CONDITION { ACTION }
CONDITION { ACTION }
END { ... }      do ACTION on lines matching CONDITION
```

so being able to get the column you want with awk is GREAT!

awk program example: sum the numbers in the 3rd column

```
action
{s += $3};
END {print s}
at the end, print the sum!
```

extract a column of text with awk

awk -F, '{print \$5}'
column separator quotes! print the 5th column
this is 99% of what I do with awk

awk program example: print every line over 80 characters

```
length($0) > 80
condition
(there's an implicit {print} as the action)
```

misc commands ❤ 17

r|wrap

adds history & ctrl support to REPLs that don't already have them (r| stands for readline)

```
$ r|wrap python
```

watch

rerun a command every 2 seconds

file

figures out what kind of file (png? pdf?) a file is

pv

"pipe viewer", gives you stats on data going through a pipe

cwl

a tiny calendar ☰

ncdu

figure out what's using all your disk space

+s

add a timestamp in front of every input line

comm

find lines 2 sorted files have in common

column

format input into columns

diff

diff 2 files. Run with -U 8 for context.

xsel / xclip

copy/paste from system clipboard. (pbcopy/pbpaste on Mac)

sort & uniq

sort sorts its input

```
$ sort names.txt
```

the default sort is alphabetical.

uniq removes duplicates

a
b
b
a
c
c

notice there are still 2 'a's! uniq only uniqueness adjacent matching lines

'sort' order	'sort -n' order
12	12
15000	48
=	=
48	96
=	=
6020	6020
=	=
96	15000

'sort -n' order	'sort -h' order
15G	45K
=	=
30M	30M
=	=
45K	15G
=	=
2006	2006

useful example:
\$ du -sh * | sort -h

sort + uniq = ♥

Pipe something to 'sort | uniq' and you'll get a deduplicated list of lines! (sort -u) does the same thing.
b
a
b
a
b
a
or sort | uniq

uniq -c

counts each line it saw.

Recipe: get the top 10 most common lines in a file:
\$ sort cats.txt
| uniq -c
| sort -n
| tail -n 10

Sed

some more sed incantations...

sed -n 12 p

print 12th line

-n suppresses output so only what you print with 'p' gets printed

sed G
(double space a file
(good for long error lines))

sed '/cat/a dog'
append 'dog' after lines containing 'cat'

sed 'i 17 panda'
insert "panda" on line 17

sed -n 5,30 p
print lines 5-30

sed 5d
delete 5th line

sed /cat/d
delete lines matching /cat/

sed -n s/cat/dog/p
only print changed lines

bash tricks

10

* **ctrl + r ***

search your history!

I use this ❤ constantly
to rerun commands

```
$ convert file.{jpg,png}  
expands to  
$ convert file.jpg file.png
```

{1..5} expands to 1 2 3 4 5

loops

```
for i in *.png
```

```
do
```

```
convert $i $i.jpg
```

```
done
```



gives the output of a command. Example:

```
$ touch file-$(date -I)
```

creates a file named file-2018-05-25

* **magical braces ***

```
$ convert file.{jpg,png}
```

expands to

{1..5} expands to 1 2 3 4 5

more keyboard shortcuts

```
ctrl+a beginning of line
```

```
ctrl+e end of line
```

```
ctrl+l clear the screen
```

+ lots more emacs
shortcuts too!

!!
expands to the last
command run
\$ sudo !!

commands that start
with a space don't go
in your history. Good if
there's a 🔒 password 🔒

top

top

a live-updating summary of
the top users of your
system's resources



let's explain some numbers in top!

15

memory

4 numbers:

total / free / used / cached

One perhaps unexpected thing:
total is not free + used!

total: free + used + cached
filesystem cache

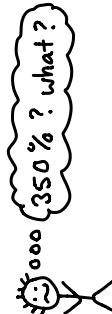
load average

3 numbers that roughly
reflect demand for your
CPUs on the system in the
last 1, 5, and 15 minutes
if it's higher than the #
of CPUs you have, that's
often bad!

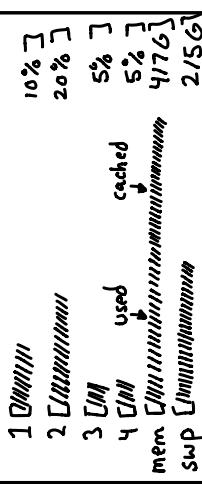
RES

this column is the "resident
set size", aka how much
RAM your process is
using.
SHR is how much of the
RES is shared with other
processes.

% CPU



this column is given as the % of
a single core. If you have
4 cores, this can go up
to 400 %!



PS

ps

ps shows which processes are running

I usually run ps like this:

```
$ ps aux
u means include username column
a+X together show all process
(ps -ef works too)
```

★ process state ★

Here's what the letters in ps's STATE column mean:

R:	running
S/D:	asleep
Z:	zombie
L:	multithreaded
+:	in the foreground

w is for wide. ps auxwww will show all the command line args for each process

e

is for environment. ps auxe will show the environment vars!

f

is for "forest" !!. ps auxf will show you an ASCII art process tree!

pstree can display a process tree, too.

you can choose which columns to show with ps (-eo ...). One cool column is 'wchan', which tells you the name of the kernel function if the process is sleeping.
try it:
\$ ps -eo user, pid, wchan, cmd

more bash tricks 11

cd -

changes to the directory you were last in

pushd & popd let you keep a stack

< ()

treat process output like a file (no more temp files!)

Example:

```
$ diff <(ls) <(ls -a)
```

ctrl+z

suspends (SIGTSTP) the running program

fg

brings backgrounded/suspended program to the foreground

bg

starts suspended program & backgrounds it (use after ctrl+z)

shellcheck

shell script linter! helps spot common mistakes.

fc

"fix command"

open the last command you ran in an editor and then run the edited version

type

tells you if something is a builtin, program, or alias
try running: \$ type time
\$ type ping
\$ type pushd
(they're all different types!)

disk usage

12

du

tells you how much disk space files / directories take up

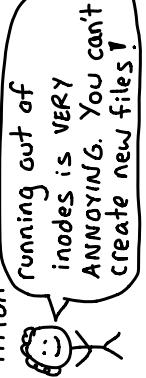
-s summary: total size of all files in a directory

-h human readable sizes

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda3	18G	6	2.5G	86%	/
udev	4.83M	4.0K	4.83M	1%	/dev
tmpfs	99M	1.4M	97M	2%	/run
/dev/sda4	167G	157G	9.9G	95%	/home

df -i

instead of % disk free, report how many inodes are used/ free on each partition



running out of inodes is VERY ANNOYING. You can't create new files!

ncdu

see what's using disk space in an interactive way

```
17.5 GiB [#####] /music  
3.2 GiB [##] /photos  
5.7 MiB [ ] /code  
2.0 MiB [ ] file.pdf
```

iostat	
get statistics about disk reads / writes	
interval to report at	'
# iostat 5	
Device:	kB_read/s kB_wrtn/s
sda	2190.21 652.87
sdb	6.00 0.00

iostat

get statistics about disk reads / writes

interval to report at

```
# iostat 5
```

Device: kB_read/s kB_wrtn/s

sda 2190.21 652.87

sdb 6.00 0.00

tar

13

The tar file format combines many files into one file.
`a.txt` `b.txt`
.tar files aren't compressed by themselves. Usually you gzip them: .tar.gz or .tgz!

-X is for extract into the current directory by default (change with -C capital C)

-C is for create lowercase makes a new tar file!

putting it together

list contents of a .tar.bz2:
\$ tar -tf file.tar.bz2
create a .tar.gz
\$ tar -czf file.tar.gz dir/
files to go in the archive

tar can compress / decompress

-Z gzip format (.gz)
-j bzip2 format (.bz2)
-J xz format (.xz)
& more! see the man page :)

-t is for list lists the contents of a tar archive

-f is for file which tar file to create or unpack