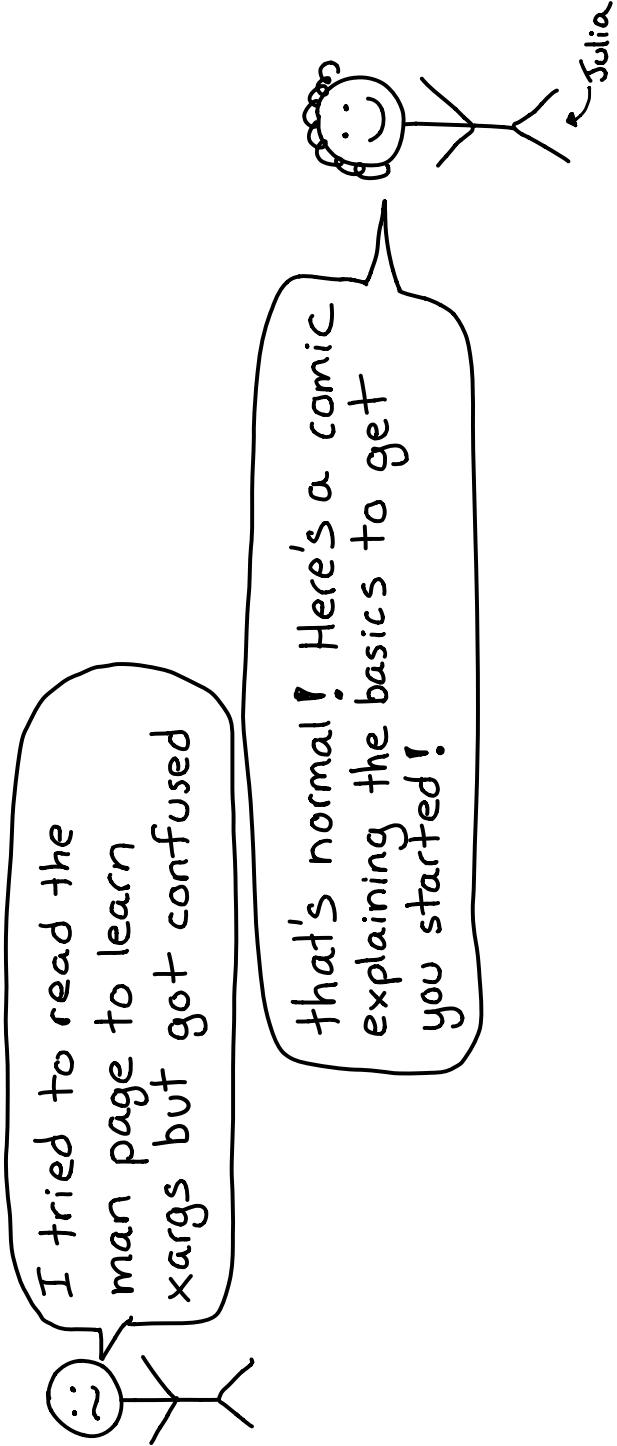




love this?
find more awesome zines at
→ wizardzines.com ←

This zine explains some of the most useful Unix command line tools in 1 page each.



Even if you've used the tool before, I might have a new trick or two for you ❤

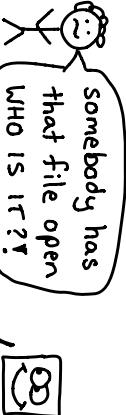
more useful tools

- make
- jq
- nohup
- disown
- cut/paste
- sponge
- xxd
- hexdump
- objdump
- strings
- screen
- tmux
- date
- entr
- seq
- join
- parallel:
 - GNU parallel
 - pigz/pixz
 - sort --parallel
- diff -U
- vipe
- imagemagick
- fish
- ranger
- chronic

lsof

lsof

stands for `list_open_files`



`ls` somebody has
that file open
WHO IS IT??

`-i`
list open network sockets
(sockets are files!)

examples:

- i -n -P → -n & -P mean "don't resolve host names / ports"
- i :8080 (also -Pni)
- i TCP
- i -s TCP:LISTEN

what `lsof` tells you
for each open file:

- pid
- file type (regular? directory?
FIFO? socket?)
- file descriptor (FD column)
- user
- filename/socket address

`lsof /some /dir`
list just the open files
in `/some /dir`

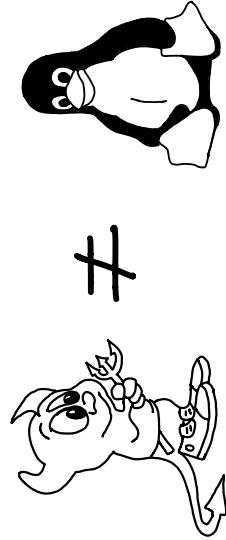
find deleted files

\$ `lsof | grep deleted`
will show you deleted files!
You can recover open deleted
files from
`/proc/PID/fd/FD`
process that opened the file

`netstat`
another way to list open
sockets on Linux is:
`netstat -tunapl`
↑
+tuna, please!
On Mac, netstat has
different args.

Table of contents

BSD#GNU	4	bash tricks	10-11	misc commands	17
grep	5	disk usage	12	head & tail	18
find	6	tar	13	less	19
xargs	7	ps	14	kill	20
awk	8	top	15	cat	21
sed	9	sort & uniq	16	lsof	22



4

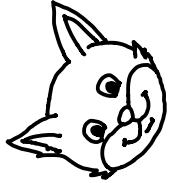
For almost all of these tools, there are at least 2 versions:

① The BSD version (on BSDs & Mac OS)

② The GNU version (on Linux)

All of the examples in this zine were tested on Linux.
Some things (like sed -i) are different on Mac.
Be careful when writing cross-platform scripts!

You can install the GNU versions on Mac with
brew install coreutils.



cat & friends

cat concatenates files

```
$ cat myfile.txt  
prints contents of myfile.txt  
$ cat *.txt  
prints all .txt files put  
together!
```

cat -n
prints out the file with
line numbers!

- 1 Once upon a midnight..
- 2 Over many a quaint..
- 3 While I nodded, nearly

cat -n

21

zcat
cats a gzipped file!

Actually just a 1-line
shell script that runs
gzip -cd, but easier
to remember.

tee

how to redirect to
a file owned by root

```
$ sudo echo "hi" >> x.txt  
this will open x.txt as your  
user, not as root, so it fails!  
$ echo "hi" | sudo tee -a x.txt  
will open x.txt as root :)
```

\$ sudo echo "hi" >> x.txt
this will open x.txt as your
user, not as root, so it fails!

```
$ echo "hi" | sudo tee -a x.txt  
will open x.txt as root :)
```

Kill

kill doesn't just kill programs

you can send ANY signal to a program with kill!

```
$ kill -SIGNAL PID
```

name or number

killall -SIGNAL NAME

signals all processes called NAME for example:

```
$ killall firefox
```

useful flags:

- w wait for all signaled processes to die
- i ask before signalling

which signal kill sends	<u>name</u>	<u>num</u>
kill	=> SIGTERM	15
kill -q	=>	SIGKILL 9
kill -KILL	=>	SIGKILL 9
kill -HUP	=>	SIGHUP 1
kill -STOP	=>	SIGSTOP 19

lists all signals	<u>kill -l</u>		
1 HUP	2 INT	3 QUIT	4 ILL
5 TRAP	6 ABRT	7 BUS	8 FPE
9 KILL	10 USR1	11 SEGV	12 USR2
13 PIPE	14 ALRM	15 TERM	16 STKFLT
17 CHLD	18 CONT	19 STOP	20 TSTP
21 TTIN	22 TTOU	23 URG	24 XCPU
25 XFSZ	26 VTIME	27 PROF	28 WINCH
29 POLL	30 PUR	31 SYS	

pgrep

prints PIDs of matching running programs

pgrep fire matches firefox
firebird
NOT bash firefox.sh

To search the whole command line (eg bash firefox.sh), use {pgrep -f}

Same as pgrep, but signals PIDs found. Example:

```
$ pkill -f firefox
```

I use pkill more than killall these days

pkill

grep

5

-E use if you want regexps like ".+" to work. otherwise you need to use ".\+" aka egrep

-V invert match: find all lines that don't match

-l only show the filenames of the files that matched

-o only print the matching part of the line instead of the whole line

-a search binaries: treat binary data like it's text instead of ignoring it!

Grep alternatives
ack **ag** **ripgrep**
(better for searching code!) ripgrep

-A Show context for your search. For example:
\$ grep -A 3 cats
will show 3 lines of context after a match

-B

-C

-F don't treat the match string as a regex
aka fgrep

find

6

find searches a directory for files
\$ find /tmp -type d, -print
↑ ↑
directory which files action to do with the files
to search

 here are my favourite find arguments!

-name/-iname
case insensitive
the filename! Example:
-name '*.txt'

-path/-ipath
search the full path!
-path '/home/*/*.go'

-type TYPE
f: regular file l: sym link
d: directory + more!

-maxdepth NUM
only descend NUM levels when searching a directory

-size 0
find empty files!
Useful to find files you created by accident

-exec COMMAND
action: run COMMAND on every file found

-print0
print null-separated filenames.
Use with xargs -0!

-delete
action: delete all files found

less

19

less is a pager
that means it lets you view (not edit) text files or piped-in text
man uses your pager (usually less) to display man pages

many vim shortcuts work in less
/ search
n/N next / prev match
j/k down / up a line
m/, mark/return to line
g/G beginning/end of file
(gg in vim)

less -r
displays bash escape codes as colours
try \$ ls --color | less -r
with -r without -r
a.+x+ a.txt → ESC[31m a.txt → ESC[31m a.txt → bold
a.txt.gz → ESC[31m a.txt.gz → ESC[31m a.txt.gz → red,bold

q quit ↵
V lowercase edit file in your \$EDITOR
arrow keys, Home/End, PgUp, PgDn work in less

F + runs a command when less starts
less +F : follow updates
less +G : start at end of file
less +20% : start 20% into file
less +/foo: search for 'foo' right away

head & tail

head

shows you the first 10 lines of a file

if you pipe a program's output to head, the program will stop after printing 10 lines (it gets sent SIGPIPE)

-c NUM

show the first/last NUM bytes of the file

\$ head -c 1k

will show the first 1024 bytes

tail

tail shows the last 10 lines!

{tail -f FILE} will follow:

print any new lines added to the end of FILE. Super useful for log files!

tail --retry

keep trying to open file if it's inaccessible

tail --pid PID

stop when process PID stops running (with -f)

-n NUM

-n NUM (either head or tail) will change the # lines shown

NUM can also be negative. Example:

\$ head -n -5 file.txt

will print all lines except the last 5

tail --follow=name

Usually tail -f will follow a file descriptor.

tail --follow=name FILENAME will keep following the same file name, even if the file descriptor changes

Xargs

7

how to replace "foo" with "bar" in all .txt files:

```
find . -name '*.txt' |  
xargs sed -i 's/foo/bar/g'
```

xargs takes whitespace separated strings from stdin and converts them into command-line arguments

echo /home/tmp | xargs ls

will run:

ls /home/tmp

how to lint every Python file in your Git repo:

```
git ls-files | grep .py |  
xargs pep8
```

more useful xargs options:

{-n 1} makes xargs run a separate process for each input
{-P} is the max number of parallel processes xargs will start

if there are spaces in your filenames "my day.txt" xargs will think it's 2 files "my" and "day.txt"

fix it like this:

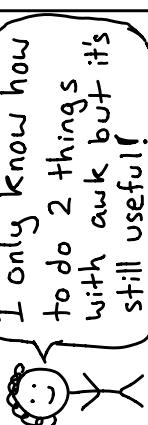
```
find . -print0 |  
xargs -0 COMMAND
```

awk

8

awk is a tiny programming language for manipulating columns of data

I only know how to do 2 things with awk but it's still useful!



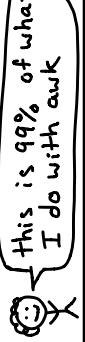
basic awk program structure:

```
BEGIN { ... }
CONDITION { ACTION }
CONDITION { ACTION }
END { ... } ↑
do ACTION on lines matching CONDITION
```

extract a column of text with awk

```
awk -F, '{print $5}'
```

↑
column separator quotes! 5th column
this is 99% of what I do with awk



so MANY unix commands print columns of text (ps! ls!)

so being able to get the column you want with awk is GREAT!

awk program example: sum the numbers in the 3rd column

```
action
{s += $3};
END {print s}
at the end, print the sum!
```

awk program example: print every line over 80 characters

```
length($0) > 80
condition
(there's an implicit {print} as the action)
```

misc commands

17

rlwrap

adds history & ctrl support to REPLs that don't already have them (rl stands for readline)

```
$ rlwrap python
```

watch

rerun a command every 2 seconds

pv

"pipe viewer", gives you stats on data going through a pipe

file

figures out what kind of file (png? pdf?) a file is

ncdu

figure out what's using all your disk space

diff

diff 2 files. Run with -U 8 for context.

xsel / xclip

copy/paste from system clipboard.
(pbcopy/pbpaste on Mac)

comm

find lines 2 sorted files have in common

Sort & uniq

sort sorts its input

```
$ sort names.txt
the default sort is alphabetical.
```

uniq removes duplicates

a
b
b => a
a
c
c
notice there are still 2 'a's! uniq only uniques adjacent matching lines

sort -n	sort -n
numeric sort	sort order
12	12
= 15000	= 30M
= 48	= 45K
= 96	= 15G
= 6020	= 200G
96	15000

useful example:
\$ du -sh * | sort -h

sort + uniq = ♥

Pipe something to 'sort | uniq' and you'll get a deduplicated list of lines! Sort -u does the same thing.
b
a | sort -u => a
b
a | sort -u => b
or sort uniq

sort -h	sort -n	sort -c
human sort	sort -n order	sort -h order
'sort -n' order	'sort -n' order	'sort -h' order
15G	45K	30M
= 30M	= 15G	= 45K
= 45K	= 200G	= 15000
= 15000	= 200G	= 30M

useful example:
\$ sort cats.txt | uniq -c | sort -n | tail -n 10

uniq -c counts each line it saw.

Recipe: get the top 10 most common lines in a file:

Sed

sed is most often used for replacing text in a file

```
$ sed s/cat/dog/g file.txt
```

can be a regular expression

sed 5d

delete 5th line

sed /cat/d

delete lines matching /cat/

sed -n 5,30P

print lines 5-30

sed s+cat/+dog/+ use + as a regex delimiter

way easier than escaping / is like s/cat\\//dog\\//!

sed -n s/cat/dog/p

only print changed lines

some more sed incantations...

sed -n 12P

print 12th line

{-n} suppresses output so only what you print with 'p' gets printed

sed G double space a file (good for long error lines)

sed '/cat/a dog' append 'dog' after lines containing 'cat'

sed '17 panda' insert "panda" on line 17

bash tricks

10

* **ctrl + r ***

search your history!

I use this **constantly** to rerun commands

* **magical braces ***

```
$ convert file.{jpg,png}  
expands to  
$ convert file.jpg file.png
```

```
{1..5} expands to 1 2 3 4 5
```

loops

```
for i in *.png  
do  
convert $i $i.jpg  
done
```

 for loops: easy & useful!

* **!!**

expands to the last command run
\$ sudo !!

commands that start with a space don't go in your history. good if there's a  password 

more keyboard shortcuts

ctrl+a	beginning of line
ctrl+e	end of line
ctrl+l	clear the screen
+ lots more emacs shortcuts too!	

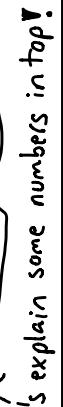
15

top

top

a live-updating summary of the top users of your system's resources

 who's using all my memory

 chrome, top abu!

let's explain some numbers in top!

load average

3 numbers that roughly reflect demand for your CPUs on the system in the last 1, 5, and 15 minutes if it's higher than the # of CPUs you have, that's often bad!

memory

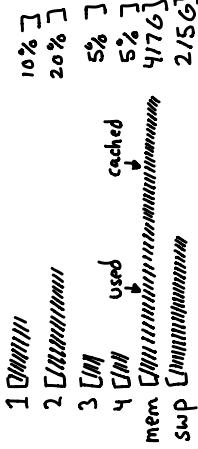
4 numbers:
total / free / used / cached
One perhaps unexpected thing: total is not free + used!
total = free + used + cached filesystem cache

15

15

htop

a prettier & more interactive version of top *



RES

this column is the "resident set size", aka how much RAM your process is using.
SHR is how much of the RES is shared with other processes

% CPU

 350%? what?

this column is given as the % of a single core. If you have 4 cores, this can go up to 400%!

PS

ps

ps shows which processes are running

I usually run ps like this:

```
$ ps aux
u means include username column
a+* together show all process
(ps -ef works too)
```

★ process state ★

Here's what the letters in ps's STATE column mean:

- R: running
- S/D: asleep
- Z: zombie
- I: multithreaded
- +: in the foreground

is for wide. ps auxww will show all the command line args for each process

e

is for environment. ps auxe will show the environment vars!

f is for "forest" !. ps auxf will show you an ASCII art process tree !

pstree can display a process tree, too.

wchan
you can choose which columns to show with ps (ps -eo ...) One cool column is 'wchan', which tells you the name of the kernel function if the process is sleeping.
try it:
\$ ps -eo user,pid,wchan,cmd

more bash tricks 11

cd -

changes to the directory you were last in

pushd & popd let you keep a stack

ctrl+z

suspends (SIGTSTP) the running program

fg

brings background/suspended program to the foreground

bg

starts suspended program & backgrounds it (use after ctrl+z)

shellcheck

shell script linter! helps spot common mistakes.

<()

process substitution

treat process output like a file (no more temp files!)

Example:

```
$ diff <(ls) <(ls -a)
```

fc

"fix command"

open the last command you ran in an editor and then run the edited version

type

tells you if something is a builtin, program, or alias

try running: \$ type time

\$ type ping

\$ type pushd

(they're all different types!)

disk usage

12

du

tells you how much disk space files / directories take up

- S summary: total size of all files in a directory
- h human readable sizes

* **df *** tells you how much free space each partition has.

-h for human-readable sizes

Filesystem	Size	Used	Avail	Use%	Mounted on
/dev/sda3	18G	6	2.5G	86%	/
udev	483M	4.0K	483M	1%	/dev
tmpfs	99M	1.4M	97M	2%	/run
/dev/sda4	167G	157G	9.9G	95%	/home

df -i

instead of % disk free, report how many inodes are used/ free on each partition

running out of inodes is VERY ANNOYING! You can't create new files!

see what's using disk space in an interactive way

17.5 GiB [#####]	/music
3.2 GiB [##]	/photos
5.7 MiB []	/code
2.0 MiB []	file.pdf

iostat

get statistics about disk reads / writes

interval to report at
iostat 5
Device: kB_read/s kB_wrtn/s
sda 2190.21 652.87
sdb 6.00 0.00

tar

13

The tar file format combines many files into one file.

- a.txt
- b.txt
- c.txt

.tar files aren't compressed by themselves. Usually you gzip them: .tar.gz or .tgz!

Usually when you use the 'tar' command, you'll run some incantation To unpack a tar.gz, use:

-tar -xzf file.tar.gz

What's xzf?
let's learn!

-X is for extract into the current directory by default (change with -C)
capital C

-C is for create lowercase makes a new tar file!

tar can compress / decompress

- Z gzip format (.gz)
- j bzip2 format (.bz2)
- J xz format (.xz)

& more! see the man page

list contents of a .tar.b2z:

- \$ tar -tf file.tar.b2z
- create a .tar.gz
- \$ tar -czf file.tar.gz dir/ files to go in the archive

-t is for list lists the contents of a tar archive

-f is for file which tar file to create or unpack