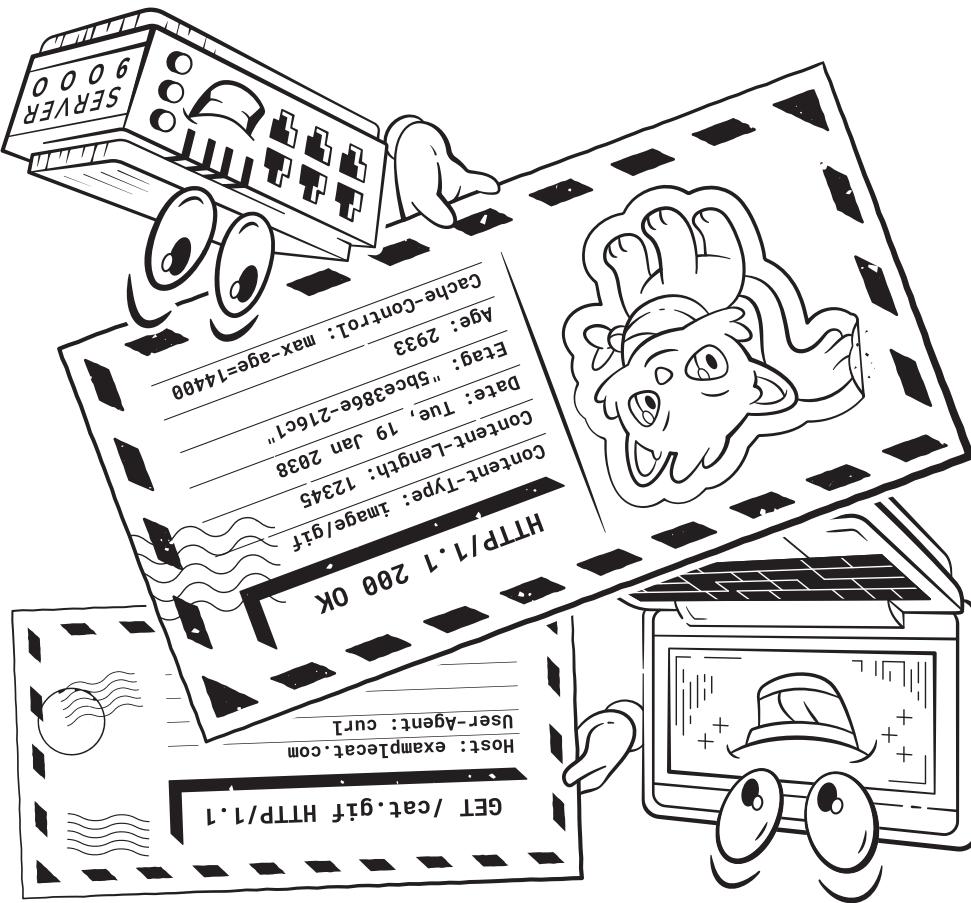


by Julia Evans



like this?  
more zines at  
wizardzines.com

# about this zine

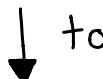
Your browser uses HTTP every time it visits a website.



This zine's goal is to take you from:



my website isn't working right  
because of some weird  
caching/cookies/CORS thing  
and I'm not sure where to start



hmm, I have a caching problem,  
I can just look at my  
request + response headers,  
consult the documentation,  
and figure out how to fix it!

## credits

Cover art: Vladimir Kašiković  
Editing: Dolly Lanuza, Kamal Marhubi  
Copy editing: Courtney Johnson  
special thanks to Marco Rogers for  
suggesting the idea of a HTTP zine

# how to learn more

## ♥ Mozilla Developer Network

<https://developer.mozilla.org>

MDN is a fantastic wiki maintained by Mozilla. It has tutorials and reference documentation for HTML, CSS, HTTP, Javascript. It's the best place to start for reference documentation.

## ♥ OWASP

<https://cheatsheetseries.owasp.org>

OWASP is an organization that publishes security best practices. If you have a question about web security, they've probably published a cheat sheet or guide to help you.

## ♥ httpstatuses.com

Nice little site that explains all the HTTP status codes.

## ♥ RFCs

<https://tools.ietf.org/html/rfcXXXX>

put RFC number here

RFCs are numbered documents (like "RFC 2631"). Every Internet protocol (like TLS or HTTP) has an RFC. These are where you go to find the Official Final Answers to technical questions you have about any internet standard. The HTTP standard is mostly documented in 6 RFCs numbered 7230 to 7235.



is the Host header  
actually required?



Yes, section 5.4  
of RFC 7230  
says so!

Don't be scared of using an RFC if you want to know for sure!

the final answer

# Table of Contents

What's HTTP?	4
How URLs work	5
What's a header?	6
Requests:	7
Anatomy of an HTTP request	8-9
Request methods (GET! POST!)	8-9
Request headers	10
Using HTTP APIs	11
Responses:	12
Anatomy of an HTTP response	12
Responses	13
How cookies work	15
Content delivery networks & caching	16-17
HTTP/2	18
Redirects	19
HTTP/2	19
Content delivery networks & caching	16-17
How cookies work	15
Status codes (200! 404!)	14
Responses headers	13
Status codes (200! 404!)	14
Content delivery networks & caching	16-17
HTTP/2	18
Redirects	19
HTTP/2	19
Same origin policy & CORS	20-21
HTTPS & certificates	20-21
Security headers	22-24
Exercises & how to learn more	26-27

Making HTTP requests with curl to real internet websites and trying different headers is my favourite way to play around with HTTP & learn.

\* curl tips

-i shows the response headers by sending a HEAD request

-I only shows the response headers -H adds a request header

curl -i https://examplecat.com/cat.txt -H "Range: bytes=8-17"

Try the Range header:

curl -i https://examplecat.com/cat.txt -H "Range: bytes=8-17"

Request (and print out!) a compressed response:

curl -i https://examplecat.com -H "Accept-Encoding: gzip" -- output -

Get a webpage in Spanish:

curl -i https://twitter.com -H "Accept-Language: es-ES"

Get redirected to another URL:

curl -i http://examplecat.com

(hint: look at the Location header!)

Guess what delivery network GitHub is using:

curl -I https://github.githubassets.com

(hint: it's in a header starting with X-)

Find out when example.com was last updated:

curl -I example.com

(hint: Last-Modified)

curl -i example.com/bananas

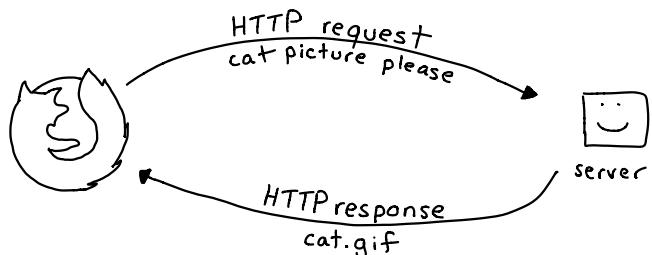
Get a 404 not found:

curl -i example.com

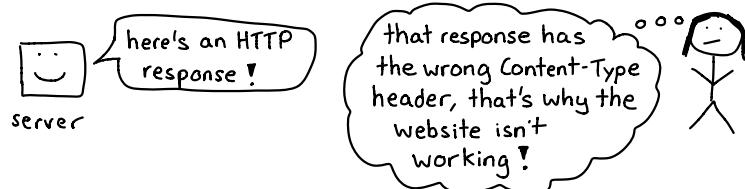
# HTTP exercises

# what's HTTP?

HTTP is the protocol (Hypertext Transfer Protocol) that's used when you visit any website in your browser.



The exciting thing about HTTP is that even though it's used for literally every website, HTTP requests and responses are easy to look at and understand:



Example of what an HTTP request and response might look like:

request	response
request line { GET / HTTP/1.1 status	HTTP/1.1 200 OK
headers { Host: examplecat.com	Cache-Control: max-age=604800
User-Agent: curl	Content-Type: text/html
Accept: */*	Etag: "1541025663+ident"
	Server: ECS (nyb/1D0B)
	Vary: Accept-Encoding
	X-Cache: HIT
	Content-Length: 1270
body { <!doctype html>	
	<title>Example Cat</title>
	...

All that text is a lot to understand, so let's get started learning what all of it means!

# security headers

These are headers your server can set. They ask the browser to protect your users' data against attackers in different ways:

**Content-Security-Policy** often called CSP

Only allow CSS / Javascript from certain domains you choose to run on your website. Helps protect against cross-site-scripting (aka XSS) attacks.

**Referrer-Policy**

Control how much information is sent to other sites in the Referrer header. Example: Referrer-Policy: no-referrer.

spelling is inconsistent with Referer header !!

**Strict-Transport-Security** often called HSTS

Require HTTPS. If you set this the client (browser) will never request a plain HTTP version of your site again. Be careful! You can't take it back!

**Expect-CT**

Certificate Transparency (CT) is a system that can help find malicious SSL certificates issued for your site. This header gives the browser a URL to use to report bad certificates to you.

**X-XSS-Protection**

Another way to protect against XSS attacks. It's not supported by all browsers, and Content-Security-Policy is more powerful.

**How URLs work**

https://examplecat.com:443/cats?color=light%20gray#banana

scheme domain port path query string fragment

Protocol to use for the request. Encrypted (https), insecure (http), or something else entirely (ftp). Where to send the request. For HTTP(s) requests, the Host header gets set to this (Host: example.com) requests to it using the Access-Control-Allow-Origin header.

If you run api.clothes.com, you can allow clothes.com to make requests to it using the Access-Control-Allow-Origin header. If you run api.clothes.com, you can allow clothes.com to make requests to it using the Access-Control-Allow-Origin header.

Defaults to 80 for HTTP and 443 for HTTPS.

Path to ask the server for. The path and the query parameters are combined in the request, like: GET /cats?color=light%20gray HTTP/1.1

Query parameters are usually used to ask for a different version of a page ("I want a light gray cat"). Example:

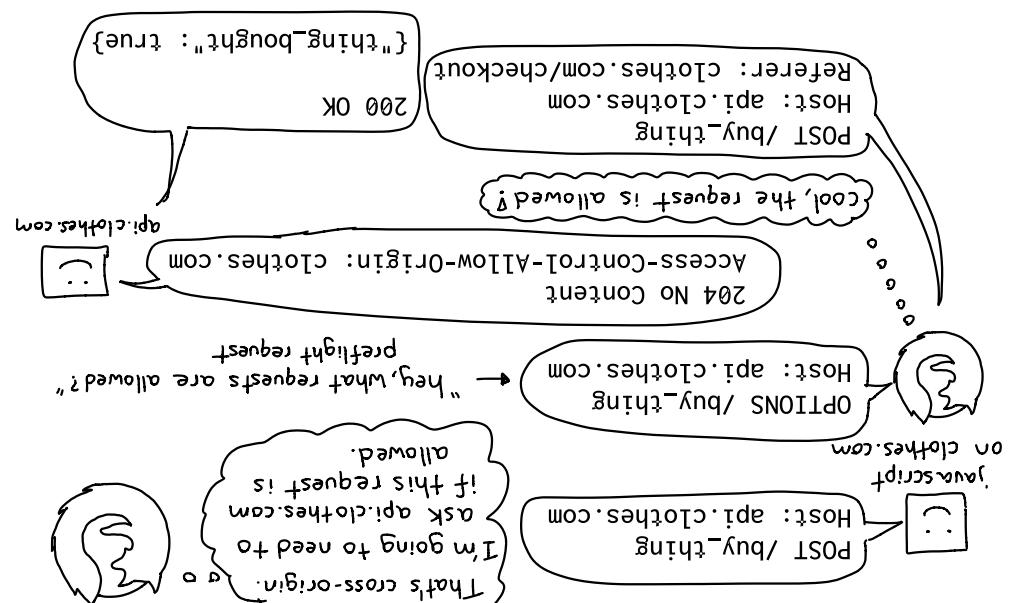
hair=short&color=black&name=mrdarci name=value separated by &

URLs aren't allowed to have certain special characters like spaces, @, etc. So to put them in a URL you need to percent encode them as % + hex representation of ASCII value.

space is %20, % is %25, etc.

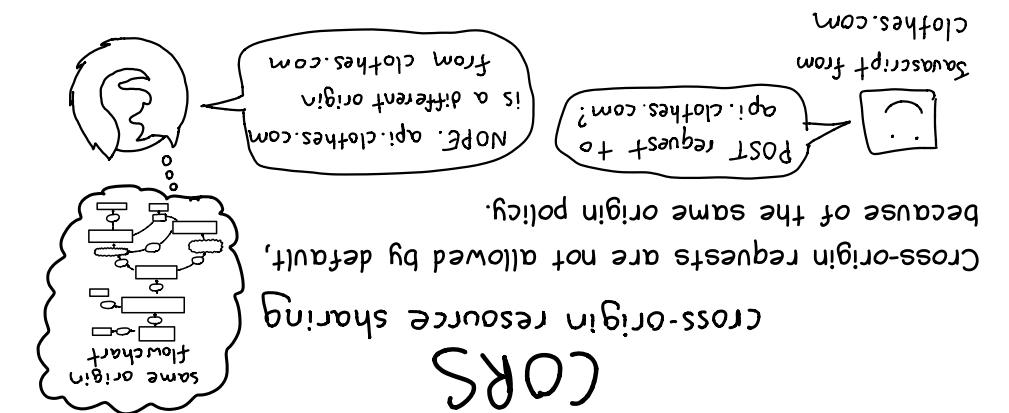
This isn't sent to the server at all. It's used either to jump to an HTML tag (<a id="banana" . . .>) or by #banana

This OPTIONS request is called a "preflight" request and it only happens for some requests, like we described in the diagram on the same-origin policy page. Most GET requests will just be sent by the browser without a preflight request first, but POST requests that send JSON need a preflight.



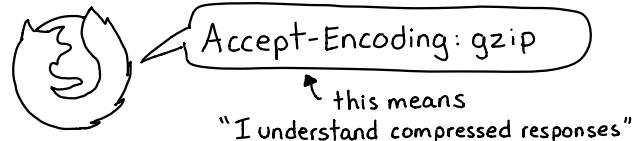
Here's what happens:

If you run api.clothes.com, you can allow clothes.com to make requests to it using the Access-Control-Allow-Origin header.

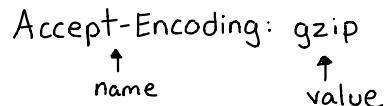


# what's a header?

Every HTTP request and response has headers. Headers are a way for the browser or server to send extra information!



Headers have a name and a value.



Header names aren't case sensitive:

totally valid ↗ aCcEPT-eNcOdInG : gzip

There are a few different kinds of headers:

Describe the body:

Content-Type: image/png Content-Encoding: gzip  
Content-Length: 12345 Content-Language: es-ES

Every Accept-  
header has a  
corresponding  
Content- header

Ask for a specific kind of response:

Accept: image/png  
Range: bytes=1-10

Accept-Encoding: gzip  
Accept-Language: es-ES

Manage caches:

ETag: "abc123"  
If-None-Match: "abc123"  
Vary: Accept-Encoding  
If-Modified-Since: 3 Aug 2019 13:00:00 GMT  
Last-Modified: 3 Feb 2018 11:00:00 GMT  
Expires: 27 Sep 2019 13:07:49 GMT  
Cache-Control: public, max-age=300

Say where the request comes from:

User-Agent: curl

Referer: https://examplecat.com

Cookies:

Set-Cookie: name=julia; HttpOnly (server → client)  
Cookie: name=julia (client → server)

6

and more!

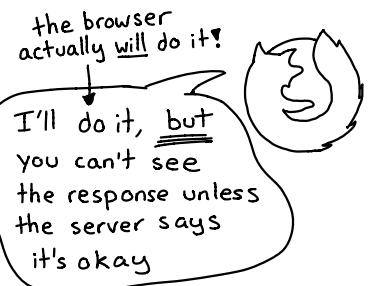
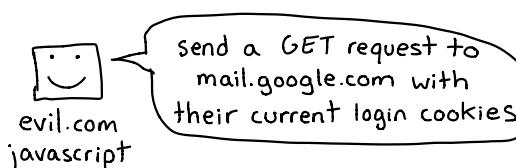
# why the same origin policy matters

Browsers work hard to make sure that evil.com can't make requests to other-website.com. But evil.com can request other-website.com from its own server. So what's the big deal?

Here are 2 reasons it's important to prevent Javascript code from making arbitrary requests from your browser:

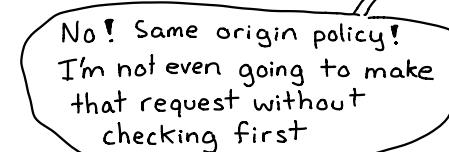
## Reason 1: cookies

Browsers often send your cookies with HTTP requests. You don't want evil.com to be able to make requests using your login cookies. They'd be logged in as you!



## Reason 2: network access

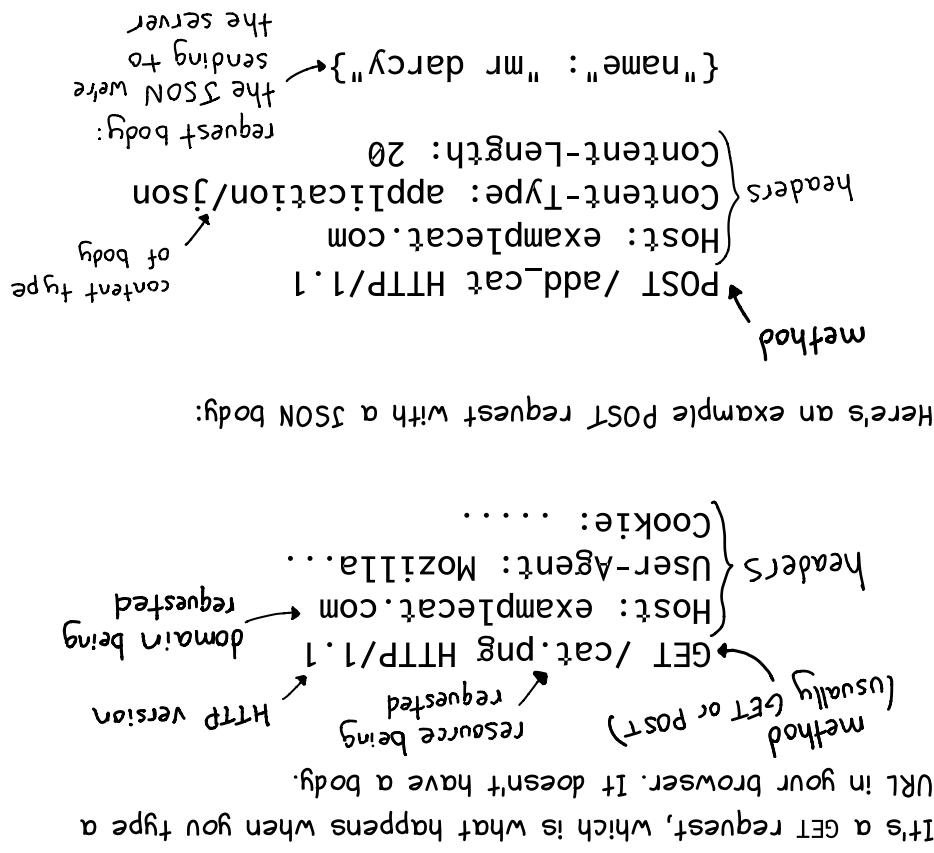
You might be on a private network (for example your company's corporate network) that evil.com doesn't have access to, but your computer does.



\* HTTP requests \*

```

graph TD
    A[HTTP requests always have:] --> B[a domain (like example.cat.com)]
    A --> C[a resource (like /cat.png)]
    B --> D[method (GET, POST, or something else)]
    B --> E[headers (extra information for the server)]
    B --> F[URL in your browser. It doesn't have a body.]
    C --> D
    C --> E
    C --> F
    D --> G[method (GET or POST)]
    D --> H[resource being requested (example/cat.png)]
    D --> I[HTTP version (HTTP/1.1)]
    E --> J[headers (User-Agent: Mozilla...)]
    E --> K[domain being requested (example.cat.com)]
    F --> L[headers (User-Agent: Mozilla...)]
    F --> M[Cookie: .....]
  
```

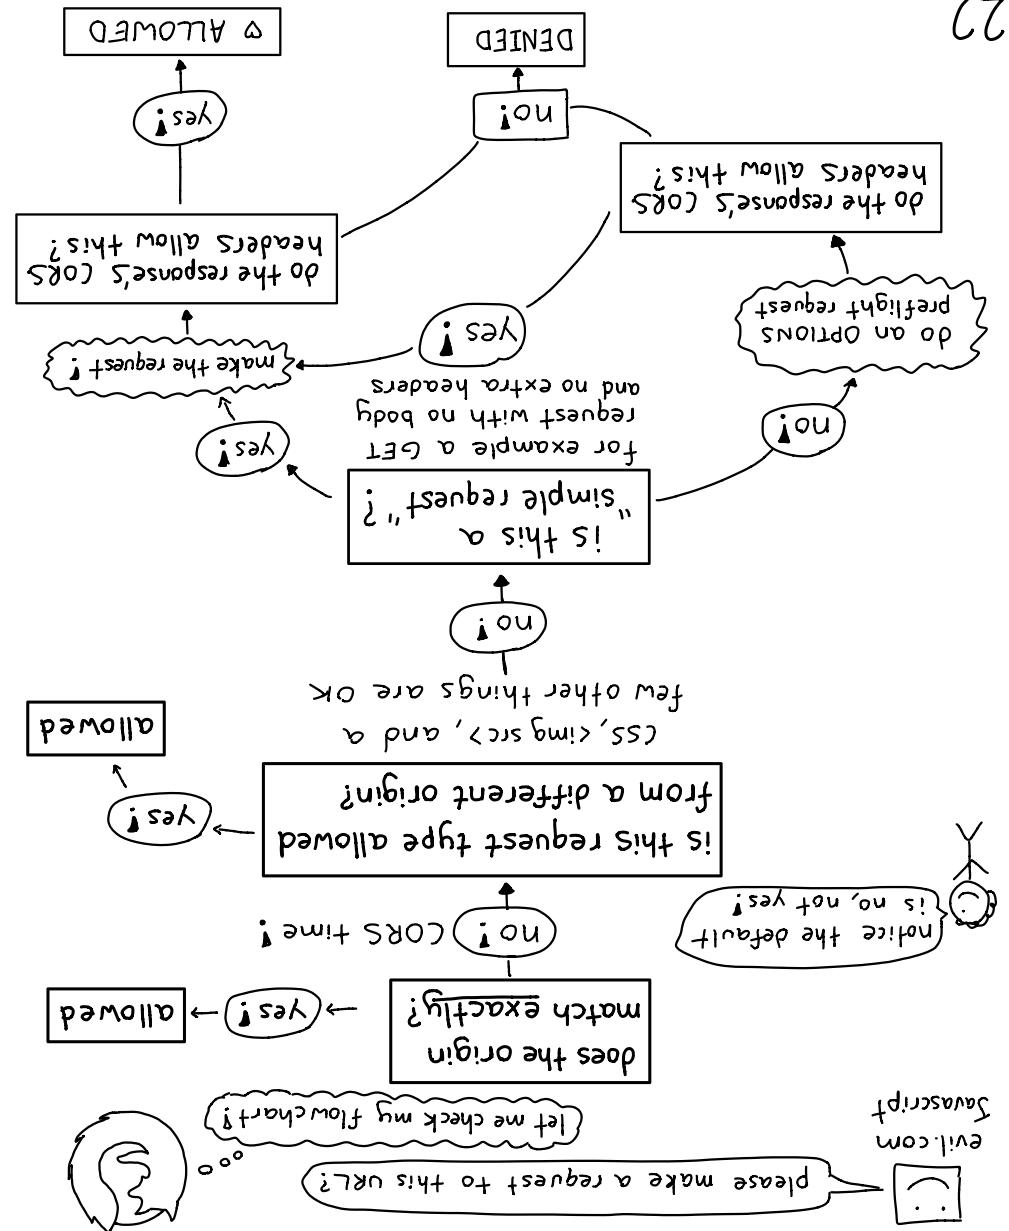


The same origin policy is one way browsers protect you from malicious JavaScript code. Here's basically how it works:

An origin is the protocol + domain including subdomains + port.  
example: <https://faby.example.com:443>

An origin is the protocol + domain including subdomains + port.

the same origin policy



# request methods

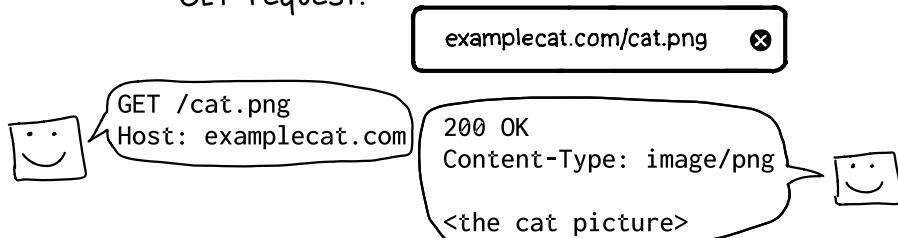
Every HTTP request has a method. It's the first thing in the first line:

↙ this means it's a *GET* request

GET /cat.png HTTP/1.1

There are 9 methods in the HTTP standard. 80% of the time you'll only use 2 (GET and POST).

**GET** When you type an URL into your browser, that's a GET request.



**POST** When you hit submit on a form, that's (usually) a POST request.



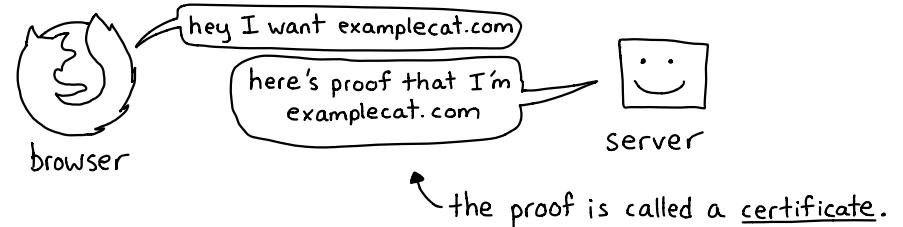
The big difference between GET and POST is that GETs are never supposed to change anything on the server.

**HEAD** Returns the same result as GET, but without the response body.



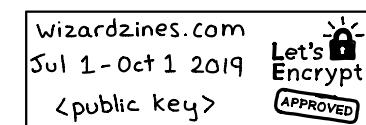
# certificates

To establish an HTTPS connection to examplecat.com, the client needs proof that the server actually is examplecat.com .



A TLS certificate has:

- a set of domains it's valid for (eg examplecat.com)
- a start and end date (example: july 1 2019 to oct 1 2019)
- a secret private key that only the server has ↗ this is the only secret part, the rest is public
- a public key to use when encrypting
- a cryptographic signature from someone trusted



The trusted entity that signs the certificate is called a ★ Certificate Authority ★ (CA) and they're responsible for verifying that you actually own the domain:



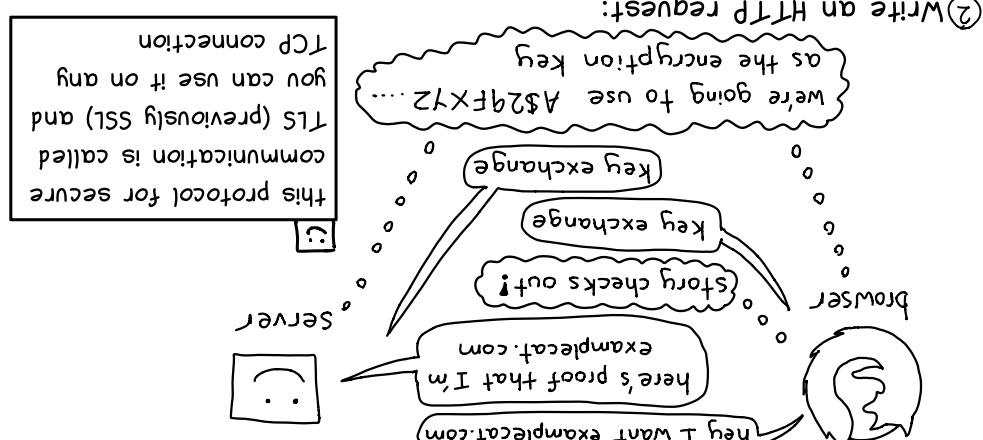
When your browser connects to examplecat.com, it validates the certificates using a list of trusted CAs installed on your computer. These CAs are called "root certificate authorities".



# HTTPS: HTTP + Secure ⚡

Here's what your browser does when it asks for `https://examplecat.com/cat.png`:

① Negotiate an encryption key (AES symmetric key) to use for this connection to `examplecat.com`. The browser and server use the same key to encrypt/decrypt content. Here's a simplified version of how picking the encryption key works:



① Negotiate an encryption key (AES symmetric key) to use in many APIs (like the Stripe API) to delete resources.

**DELETE** Used in many APIs (like the Stripe API) to delete resources.

`DELETE /v1/customers/cus_12345`

"Delete this customer, please!"

"Deleted ▶"

"200 OK"

If also tells you which methods are available.  
Page 24 has more about CORS.

**OPTIONS** OPTIONS is mostly used for CORS requests.

**PUT** Used in some APIs (like the S3 API) to create or update resources. PUT /cat/1234 lets you GET /cat/1234 later.

**TRACE**

I've never seen a server that supports this, you probably don't need to know about it.

resource ("just change this file!").

Used in some APIs for partial updates to a resource ("just change this file!").

**CONNECT** Different from all the others: instead of making a request directly to a server, it asks for a proxy to open a connection.

If you set the HTTPS\_PROXY environment variable to a proxy server, many HTTP libraries will use this protocol to proxy your requests.



③ Encrypt the HTTP request with AES & send it to `examplecat.com`:

② Write an HTTP request:

Host: examplecat.com  
GET /cat.png HTTP/1.1  
User-Agent: Mozilla/5.0

...  
ah I see,  
\$A79bbc-a  
~99bf..

④ Receive encrypted HTTP response:

Content-Type: image/png  
200 OK  
nice, that means:

BXF a56□gxx ...

# request headers

These are the most important request headers:

## Host

The domain.  
The only required header.

Host: examplecat.com User-Agent: curl 7.0.2 Referer: https://examplecat.com

## User-Agent

name + version of your browser and OS

↑ yes, it's misspelled!

## Referer

website that linked or included the resource

## Authorization

eg a password or API token

base64 encoded user:password

Authorization: Basic YXZ

## Cookie

Send cookies the server sent earlier.  
Keeps you logged in.

Cookie: user=b0rk

## Range

lets you continue downloads ("get bytes 100 - 200")  
Range: bytes=100-200

## Cache-Control

"max-age = 60" means cached responses must be less than 60 seconds old

## If-Modified-Since

only send if resource was modified after this time

If-Modified-Since: Wed, 21 Oct... If-None-Match: "e7ddac"

## If-None-Match

only send if the ETag doesn't match those listed

## Accept

MIME type you want the response to be

Accept: image/png

## Accept-Encoding

set this to "gzip" and you'll probably get a compressed response

Accept-Encoding: gzip

## Accept-Language

set this to "fr-CA" and you might get a response in French

Accept-Language: fr-CA

## Content-Type

MIME type of request body, e.g. "application/json"

## Content-Encoding

will be "gzip" if the request body is gzipped

## Connection

"close" or "keep-alive". Whether to keep the TCP connection open

# HTTP/2

HTTP/2 is a new version of HTTP.

Here's what you need to know:

## ★ A lot isn't changing

All the methods, status codes, request/response bodies, and headers mean exactly the same thing in HTTP/2.

before (HTTP/1.1)

method: GET

path: /cat.gif

headers:

- Host: examplecat.com

- User-Agent: curl

after (HTTP/2)

method: GET

path: /cat.gif

headers: authority: examplecat.com

- User-Agent: curl

## ★ HTTP/2 is faster

Even though the data sent is the same, the way HTTP/2 sends it is different. The main differences are:

- It's a binary format (it's harder to tcpdump traffic and debug)
- Headers are compressed
- Multiple requests can be sent on the same connection at a time

before (HTTP/1.1)

→ request 1

response 1 ←

→ request 2

response 2 ←

after (HTTP/2)

→ request 1

→ request 2

out of order is OK

{ response 2 ←

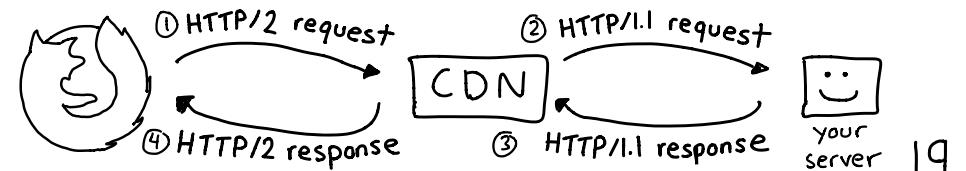
response 1 ← }

} one TCP connection

All these changes together mean that HTTP/2 requests often take less time than the same HTTP/1.1 requests.

## ★ Sometimes you can switch to it easily

A lot of software (CDNs, nginx) let clients connect with HTTP/2 even if your server still only supports HTTP/1.1.



## Using HTTP APIs

- Lots of services (Twitter! Tumblr! Google!) let you use them by sending them HTTP requests. If an HTTP API doesn't come with a client library, don't be scared! You can just make the HTTP requests yourself. Here's what you need to remember:
  - Set the right Content-Type header
  - Often you'll be sending a POST request with a body, and that means you need a Content-Type header that matches the body.

- \* application/json → JSON;
- \* application/x-www-form-urlencoded → Same as what a HTML form does

a common error is to try to send POST data as one content type (like JSON) when it's actually another like application/x-www-form-urlencoded

Identify yourself  
Most HTTP APIs require a secret API key so they know who you are.  
Here's how that looks for the Twilio API:

```
api.twitter.com/2010-04-01/Accounts/ACCOUNT_ID/Messages.json  
Content-Type: application/json  
ACCOUNT_ID: AUTH_TOKEN  
-u sends the username/password  
"from": "+15141234567",  
"to": "+15141234567",  
"body": "a text message"  
in the Authorization header  
this sends a POST request
```

301 Moved Permanently! redirects are PERMANENT: after a browser sees one once, it'll always use example.cat.com/cat.png when someone types example.cat.com/dog.png, forever. You can't take it back and decide to not to redirect. If you're not sure you want to redirect your site for eternity, use 302 Found to redirect instead.

Warning!

The `Location` header tells the browser what new URL to use. The new URL doesn't have to be on the same domain: `examplecat.com/panda` can redirect to `pandas.com`. Setting up redirects is a great thing to do if you move your site to a new domain!

Here's what's going on behind the scenes:

but you end up at a slightly different URL, like this:  
oooh, where did the cat come from?  
I didn't type that!

[examplecat.com/cat.png](http://examplecat.com/cat.png)

 examplecat.com/dog.png

Sometimes you type a URL into your browser like this:

redirects

# anatomy of an HTTP response

HTTP responses have:

- a status code (200 OK! 404 not found!)
- headers
- a body (HTML, an image, JSON, etc)

Here's the HTTP response from examplecat.com/cat.txt:

```

HTTP/1.1 200 OK status
Accept-Ranges: bytes
Cache-Control: public, max-age=0
Content-Length: 33
Content-Type: text/plain; charset=UTF-8
Date: Mon, 09 Sep 2019 01:57:35 GMT
Etag: "ac5affa59f554a1440043537ae973790-ssl"
Strict-Transport-Security: max-age=31536000
Age: 0
Server: Netlify

\ ^ cat! ^
) ( ' ) 
( / ) 
\(_)
    
```

} status code  
headers  
body

There are a few kinds of response headers:

- when the resource was sent/modified:

```

Date: Mon, 09 Sep 2019 01:57:35 GMT
Last-Modified: 3 Feb 2017 13:00:00 GMT
    
```

- about the response body:

```

Content-Language: en-US Content-Type: text/plain; charset=UTF-8
Content-Length: 33 Content-Encoding: gzip
    
```

- caching:

```

ETag: "ac5affa..." Age: 255
Vary: Accept-Encoding Cache-Control: public, max-age=0
    
```

- security: (see page 25)

```

X-Frame-Options: DENY Strict-Transport-Security: max-age=31536000
X-XSS-Protection: 1 Content-Security-Policy: default-src https:
    
```

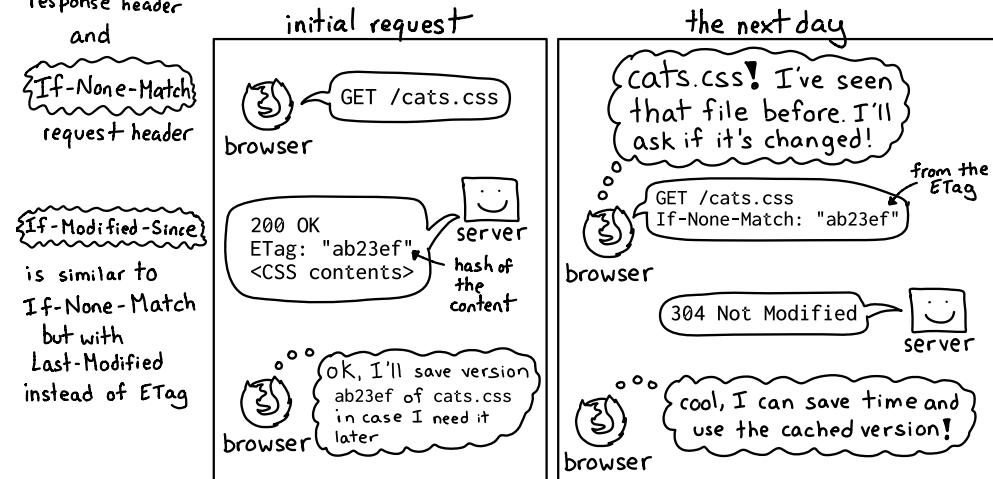
- and more:

```

Connection: keep-alive Accept-Ranges: bytes
Via: nginx
Set-Cookie: cat=darcy; HttpOnly; expires=27-Feb-2020 13:18:57 GMT;
    
```

# caching headers

These 3 headers let the browser avoid downloading an unchanged file a second time.



Vary  
response header

Sometimes the same URL can have multiple versions (Spanish, compressed or not, etc). Caches categorize the versions by request header like this:

Accept-Language	Accept-Encoding	content
en-US	-	hello
es-ES	-	hola
en-US	gzip	f\$xx99&ef^.. (compressed gibberish)

The Vary header tells the cache which request headers should be the columns of this table.

Cache-Control

request AND response header

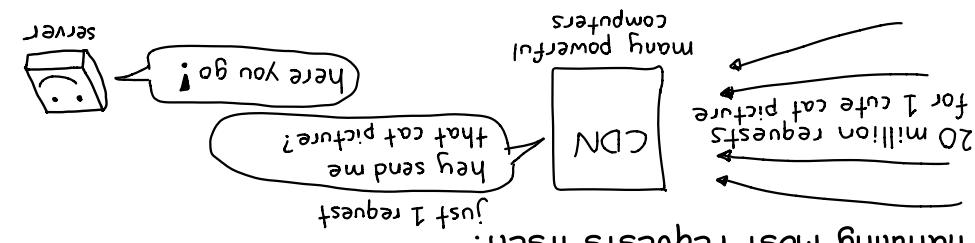
Used by both clients and servers to control caching behaviour. For example:

Cache-Control: max-age=9999999999  
from the server asks the CDN or browser to cache the thing for a long time.

## Content delivery networks

In 2004, if your website suddenly got popular, often the webserver wouldn't be able to handle all the requests.

A CDN (content delivery network) can make your site faster and save you money by caching your site and handling most requests itself.



This is great but caching can cause problems too!

Today, there are many free or cheap CDN services available, which means if your site gets popular suddenly you can easily keep it running!

I updated my site yesterday but people are still seeing the old site!

French users are seeing the English site?? Why?

Vary header

Next, we'll explain the HTTP headers your CDN or browser uses to decide how to do caching.

13

Whether Range request header is supported for this resource

Accept-Ranges

Content-Encoding: gzip

Whether body is compressed

Content-Encoding

Content-Length: 33

Length of body in bytes

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Connection

Close" or "keep-alive

Vary

Last-Modified

Date

Age

When response was sent

Seconds response was sent

Has been cached

Date: Mon, 09 Sep 2019...

Age: 355

Version of response body

Etag: "ac5afffa.."

Cache-control: max-age=300

Various caching settings

Cache-control

Request headers will that response will

Vary based on

HTTP headers

Expires

The response is stale after this time.

Proxy servers added by via: negotia

TCP connection open whether to re-requested and should be re-requested after this time.

Set-Cookie: name=value; HttpOnly

Sets a cookie.

Content-Type

MIME type of body

Content-Type

Allows CORS requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

Content-Length

Length of body in bytes

Content-Encoding

Whether body is compressed

Content-Encoding: gzip

Accept-Ranges

Whether Range request header is supported for this resource

Content-Type

Allow cross-origin requests. These called CORS headers.

Access-Control-\*

# HTTP status codes

Every HTTP response has a ★status code★



There are 50ish status codes, and these are the most common ones in real life:

200 OK

} 2xx's mean  
★ success ★

301 Moved Permanently

} 3xx's aren't errors, just redirects to somewhere else

302 Found

temporary redirect

304 Not Modified

the client already has the latest version, "redirect" to that

400 Bad Request

} 4xx errors generally mean the client made an invalid request

403 Forbidden

API key/OAuth/something needed

404 Not Found

we all know this one :)

429 Too Many Requests

you're being rate limited

} 5xx errors generally mean something's wrong with the server

500 Internal Server Error

the server code has an error

503 Service Unavailable

could mean nginx (or whatever proxy)

couldn't connect to the server

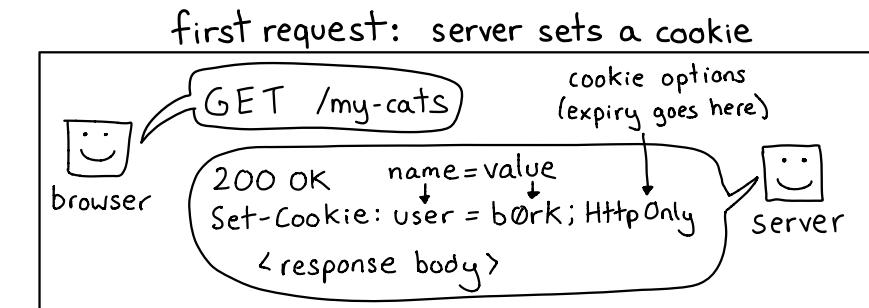
504 Gateway Timeout

the server was too slow to respond

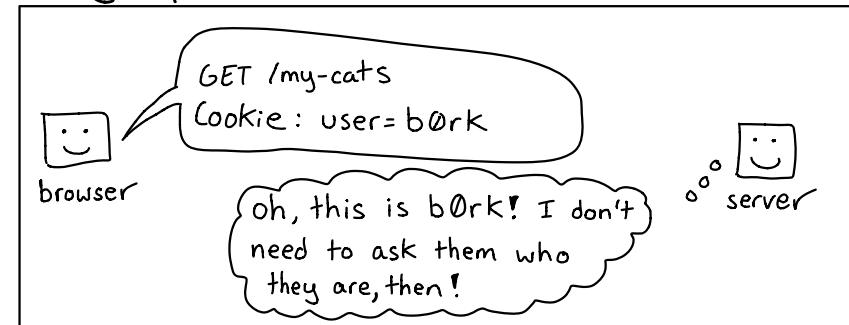
# how cookies work

Cookies are a way for a server to store a little bit of information in your browser.

They're set with the Set-Cookie response header, like this:



Every request after: browser sends the cookie back



Cookies are used by many websites to keep you logged in. Instead of user=b0rk they'll set a cookie like sessionid=long-incomprehensible-id. This is important because if they just set a simple cookie like user=b0rk, anyone could pretend to be b0rk by setting that cookie!

Designing a secure login system with cookies is quite difficult. To learn more about it, google "OWASP Session Management Cheat Sheet".